

# Visualising High Dimensional Data with t-SNE

Harri Edwards

# t-SNE???

- ▶ t-SNE is a technique for transforming high-dimensional data

$$x_1, \dots, x_k \in \mathbb{R}^{D>3}$$

to lower dimensional ‘map points’

$$y_1, \dots, y_k \in \mathbb{R}^{d \leq 3}$$

for the purpose of visualisation.

# t-SNE???

- ▶ t-SNE is a technique for transforming high-dimensional data

$$x_1, \dots, x_k \in \mathbb{R}^{D>3}$$

to lower dimensional 'map points'

$$y_1, \dots, y_k \in \mathbb{R}^{d \leq 3}$$

for the purpose of visualisation.

- ▶ Desideratum for a mapping is to *preserve the neighbourhoods* as much as possible, *i.e.* if  $x_i, x_j$  are 'close' then  $y_i, y_j$  should be 'close' and vice-versa.

# t-SNE???

- ▶ t-SNE is a technique for transforming high-dimensional data

$$x_1, \dots, x_k \in \mathbb{R}^{D>3}$$

to lower dimensional ‘map points’

$$y_1, \dots, y_k \in \mathbb{R}^{d \leq 3}$$

for the purpose of visualisation.

- ▶ Desideratum for a mapping is to *preserve the neighbourhoods* as much as possible, *i.e.* if  $x_i, x_j$  are ‘close’ then  $y_i, y_j$  should be ‘close’ and vice-versa.
- ▶ t-SNE comes from the paper ‘Visualizing Data using t-SNE’ published in 2008 by Laurens van der Maaten and Geoff Hinton (1127 citations!)

# SNE

- ▶ t-SNE is a modification of an earlier technique called 'Stochastic Neighbour Embedding' from Hinton and Roweis (2002).

# SNE

- ▶ t-SNE is a modification of an earlier technique called 'Stochastic Neighbour Embedding' from Hinton and Roweis (2002).

The basic idea of SNE is very simple: for each data point  $x_i$ , we specify a probability distribution over its possible neighbours  $x_j$  with  $j \neq i$

$$p_{i|j} = \frac{\exp(-\|x_i - x_j\|^2 / 2\sigma_i^2)}{\sum_{j \neq i} \exp(-\|x_i - x_j\|^2 / 2\sigma_i^2)} \quad (1)$$

# SNE

- ▶ t-SNE is a modification of an earlier technique called 'Stochastic Neighbour Embedding' from Hinton and Roweis (2002).

The basic idea of SNE is very simple: for each data point  $x_i$ , we specify a probability distribution over its possible neighbours  $x_j$  with  $j \neq i$

$$p_{i|j} = \frac{\exp(-\|x_i - x_j\|^2 / 2\sigma_i^2)}{\sum_{j \neq i} \exp(-\|x_i - x_j\|^2 / 2\sigma_i^2)} \quad (1)$$

Now we do the same for the map points:

$$q_{i|j} = \frac{\exp(-\|y_i - y_j\|^2)}{\sum_{j \neq i} \exp(-\|y_i - y_j\|^2)} \quad (2)$$

## SNE: Cost

So we have a probability distribution  $p$  over the neighbourhood structure of the  $x$ 's and a probability distribution  $q$  over the neighbourhood structure of the map points.



## SNE: Cost

So we have a probability distribution  $p$  over the neighbourhood structure of the  $x$ 's and a probability distribution  $q$  over the neighbourhood structure of the map points.

Making a visualisation that preserves the neighbourhood structure then consists in learning a  $q$  that matches  $p$ . The parameters to be learned are the  $y_i$ , and we measure the goodness of fit by the KL-divergence

$$C = \sum_i KL(p_{(\cdot|i)} \| q_{(\cdot|i)}) \quad (3)$$

## SNE: Cost

So we have a probability distribution  $p$  over the neighbourhood structure of the  $x$ 's and a probability distribution  $q$  over the neighbourhood structure of the map points.

Making a visualisation that preserves the neighbourhood structure then consists in learning a  $q$  that matches  $p$ . The parameters to be learned are the  $y_i$ , and we measure the goodness of fit by the KL-divergence

$$C = \sum_i KL(p_{(\cdot|i)} \| q_{(\cdot|i)}) \quad (3)$$

that is

$$C = \sum_i \sum_j p_{j|i} \log \frac{p_{j|i}}{q_{j|i}} \quad (4)$$

## SNE: Gradient

The gradient of the cost function has a nice and simple form:

$$\frac{\partial C}{\partial y_i} = 2 \sum_j (p_{j|i} - q_{j|i} + p_{i|j} - q_{i|j}) (y_i - y_j) \quad (5)$$

## SNE: Gradient

The gradient of the cost function has a nice and simple form:

$$\frac{\partial C}{\partial y_i} = 2 \sum_j (p_{j|i} - q_{j|i} + p_{i|j} - q_{i|j}) (y_i - y_j) \quad (5)$$

The optimisation of  $C$  involves initialising the map points from a spherical Gaussian, using gradient descent with momentum, and adding noise to the map points on each iteration with annealing variance.

# The Crowding Problem

'High dimensional space is big. You just won't believe how vastly, hugely, mind-bogglingly big it is. I mean, you may think it's a long way down the road to the chemist's, but that's just peanuts to space.' (Euclid)

# The Crowding Problem

‘High dimensional space is big. You just won’t believe how vastly, hugely, mind-bogglingly big it is. I mean, you may think it’s a long way down the road to the chemist’s, but that’s just peanuts to space.’ (Euclid)

- ▶ The volume of a ball with radius  $r$  in dimension  $k$  scales as  $r^k$ . This means that if you want to faithfully represent distances in 2 dimensions then you are going to have to spread out, *a lot*.

# The Crowding Problem

‘High dimensional space is big. You just won’t believe how vastly, hugely, mind-bogglingly big it is. I mean, you may think it’s a long way down the road to the chemist’s, but that’s just peanuts to space.’ (Euclid)

- ▶ The volume of a ball with radius  $r$  in dimension  $k$  scales as  $r^k$ . This means that if you want to faithfully represent distances in 2 dimensions then you are going to have to spread out, *a lot*.
- ▶ Since we used a Gaussian distribution for our  $q_{i|j}$  probabilities, the probability of being neighbours decays square-exponentially.
- ▶ Spreading out (*a lot*) + square-exponential decay = very small  $q_{i|j}$  for medium  $p_{i|j}$ , this makes the KL-divergence very sad.

# The Crowding Problem

'High dimensional space is big. You just won't believe how vastly, hugely, mind-bogglingly big it is. I mean, you may think it's a long way down the road to the chemist's, but that's just peanuts to space.' (Euclid)

- ▶ The volume of a ball with radius  $r$  in dimension  $k$  scales as  $r^k$ . This means that if you want to faithfully represent distances in 2 dimensions then you are going to have to spread out, *a lot*.
- ▶ Since we used a Gaussian distribution for our  $q_{i|j}$  probabilities, the probability of being neighbours decays square-exponentially.
- ▶ Spreading out (*a lot*) + square-exponential decay = very small  $q_{i|j}$  for medium  $p_{i|j}$ , this makes the KL-divergence very sad.
- ▶ SNE's solution to this problem? Compress all the map points to a blob so everyone is friends with everyone (in particular nobody is not friends on the map who is friends in  $x$ -space.)



# Fat Tails to the Rescue

An alternative to the group-hug solution is to use  $q$  distributions that are less needy, and more comfortable with long-distance relationships.

## Fat Tails to the Rescue

An alternative to the group-hug solution is to use  $q$  distributions that are less needy, and more comfortable with long-distance relationships. We are talking of course about the Student  $t$ -distribution with one degree of freedom:

$$q_{i,j} = \frac{(1 + \|y_i - y_j\|^2)^{-1}}{\sum_{k \neq l} (1 + \|y_k - y_l\|^2)^{-1}} \quad (6)$$

## Fat Tails to the Rescue

An alternative to the group-hug solution is to use  $q$  distributions that are less needy, and more comfortable with long-distance relationships. We are talking of course about the Student  $t$ -distribution with one degree of freedom:

$$q_{i,j} = \frac{(1 + \|y_i - y_j\|^2)^{-1}}{\sum_{k \neq l} (1 + \|y_k - y_l\|^2)^{-1}} \quad (6)$$

This behaves like an inverse-square law for larger distances, meaning that the probabilities are somewhat invariant to changes in scale for further apart map points.

## Fat Tails to the Rescue


An alternative to the group-hug solution is to use  $q$  distributions that are less needy, and more comfortable with long-distance relationships. We are talking of course about the Student  $t$ -distribution with one degree of freedom:

$$q_{i,j} = \frac{(1 + \|y_i - y_j\|^2)^{-1}}{\sum_{k \neq l} (1 + \|y_k - y_l\|^2)^{-1}} \quad (6)$$


This behaves like an inverse-square law for larger distances, meaning that the probabilities are somewhat invariant to changes in scale for further apart map points.

Also note that they use a symmetrised  $q_{i,j}$  as opposed to  $q_{i|j}$ , this makes the computations simpler, but is not conceptually important.

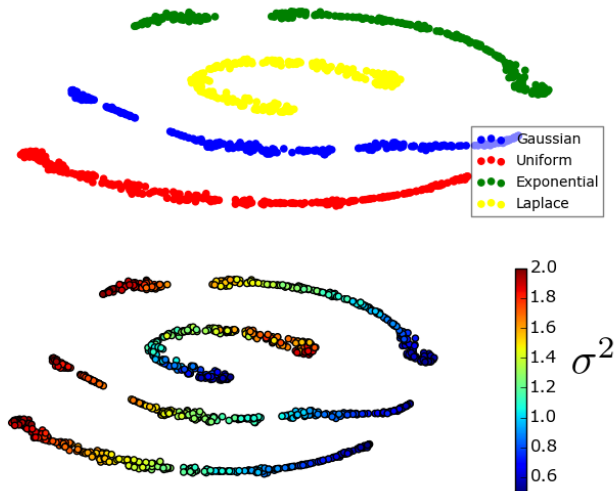
# Pictures!!!: MNIST



# Pictures!!!: Subreddits



# Pictures!!!: Neural Statistician



# Acknowledgements



## EPSRC

Engineering and Physical Sciences  
Research Council



# THE UNIVERSITY *of* EDINBURGH



THE UNIVERSITY *of* EDINBURGH

**informatics**